


COMP 7640

Written Assignment #3

Due: 11:59 PM April 24 (Friday), 2025


Question 1. Given an extendible hashing index as shown in the figure below, where each bucket can contain 4 data entries,

- (1) If we are given a query with the condition $\text{key}=11$, how many I/Os are needed for this query? **(6 marks)**
- (2) Draw the resulting index after inserting 2^* and then 8^* . **(16 marks)**



Solution:

- (1) 2 I/Os. 1 I/O for reading Bucket D (**3 marks**). 1 I/O for reading the data page. (**3 marks**)
- (2) Below see the updated index



Question 2. We are given two relations S and E, which contain 10,000 records and 1,000 records, respectively. Each disk page can hold 100 S records and 5 E records, respectively. Consider the relational operation $S \bowtie_{S.sid=E.sid} E$ and estimate the I/O costs of the following access paths.

- (a) What is the lowest cost of using a page-oriented nested loop join? **(10 marks)**
- (b) What is the lowest cost of using a simple-nested loop join? **(12 marks)**

[**Remark:** There is no I/O cost for writing the result back to the disk.]

Solution:

#Pages for S: $10,000/100 = 100$ pages **(4 marks)**

#Pages for E: $1,000/5 = 200$ pages **(4 marks)**

(a) $100 + 100 * 200 = 20,100$ I/Os **(4 marks)**

(b) When S is the outer relation: $100 + 10,000 * 200 = 2,000,100$ I/Os **(4 marks)**

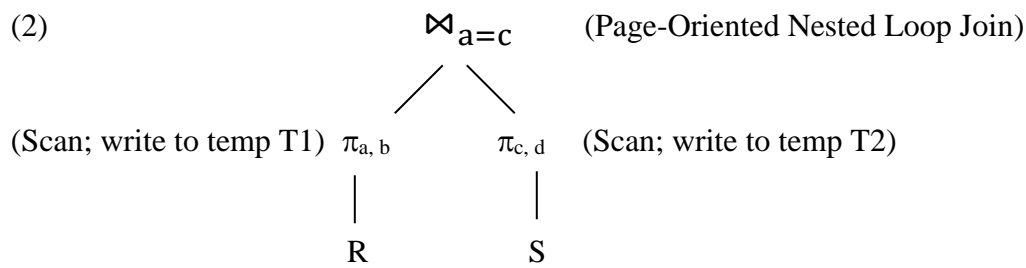
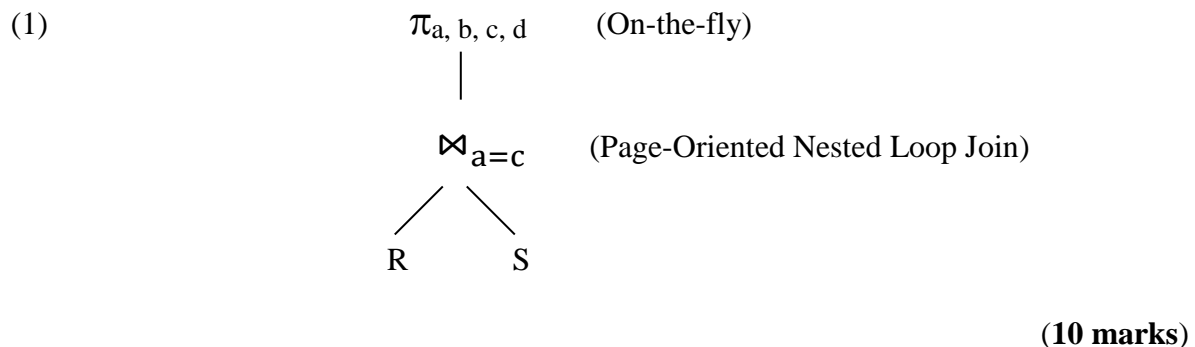
When E is the outer relation: $200 + 1,000 * 100 = 100,200$ I/Os **(4 marks)**

So the lowest I/O cost is 100,200 I/Os. **(2 marks)**

Question 3. Consider the query $\pi_{a,b,c,d} (R \bowtie_{a=c} S)$ with the following information:

- a) Relation R(a, b, e) (a is the key) has 10 pages, and each R record needs 300 bytes;
- b) Relation S(c, d, f, g) (c is the key) has 100 pages, and each S record needs 500 bytes;
- c) Attributes a, b, c, and d take 100, 50, 100, and 150 bytes, respectively;
- d) The page size is 1024 bytes;
- e) Each S record joins with exactly one R record.

Estimate the total I/O cost for the following two query evaluation plans:



(22 marks)

Solution:

(1) Cost of join = $10 + 10 * 100 = 1010$ page I/Os (5 marks)

Total cost = Cost of join = 1010 page I/Os (5 marks)

(2)

Cost of $\pi_{a,b}(R)$

= Cost to scan R + cost to write T1

= 10 pages + size(T1)

= 10 pages + (150 bytes/300 bytes) * 10 pages

= 10 pages + 5 pages

= 15 pages (7 marks)

Cost of $\pi_{c,d}(S)$

= Cost to scan S + cost to write T2

= 100 pages + size(T2)

= 100 pages + (250 bytes/500 bytes) * 100 pages

= 100 pages + 50 pages

= 150 pages (7 marks)

Cost of join = $5 + (5 * 50) = 255$ pages (4 marks)

Total cost = Cost of projection + Cost of join = $165 + 255 = 420$ I/Os (4 marks)

(The following is also correct :

#records in R = $10 * \lfloor 1024 / 300 \rfloor = 30$

#T1-records per page = $\lfloor 1024 / 150 \rfloor = 6$

Size(T1) = $30 / 6 = 5$ pages

Cost of $\pi_{a,b}(R)$

= Cost to scan R + cost to write T1

= 10 pages + 5 pages

= 15 pages (7 marks)

#records in S = $100 * \lfloor 1024 / 500 \rfloor = 200$

#T2-records per page = $\lfloor 1024 / 250 \rfloor = 4$

Size(T2) = $200 / 4 = 50$ pages

Cost of $\pi_{c,d}(S)$

= Cost to scan S + cost to write T2

= 100 pages + 50 pages

= 150 pages (7 marks)

Cost of join = $5 + (5 * 50) = 255$ pages (4 marks)

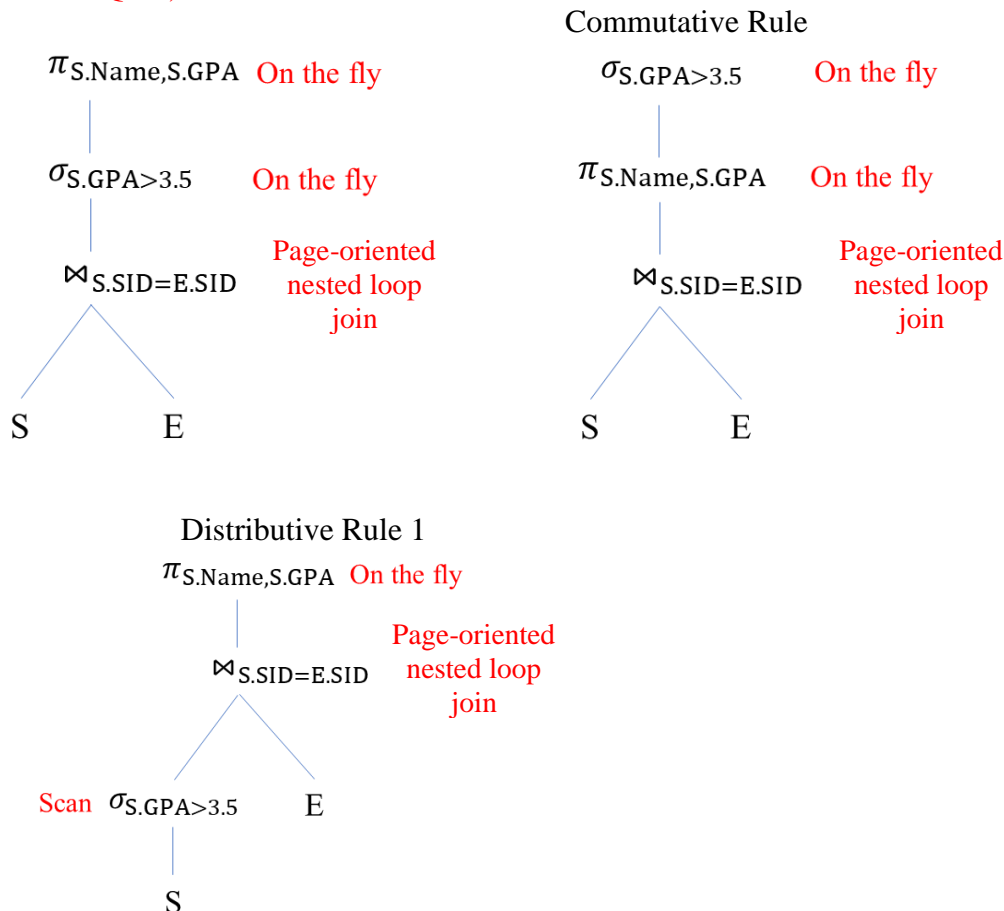
Total cost = Cost of projection + Cost of join = $165 + 255 = 420$ I/Os (4 marks)

Question 4. Given two relations S(SID, Name, GPA, age) and E(CID, SID, Date), and the following SQL query:

```
SELECT S.Name, S.GPA
FROM S, E
WHERE S.GPA>3.5 AND S.SID=E.SID
```

Draw three query evaluation plans for this SQL query. (24 marks)

Solution: Other appropriate access paths are also correct. (No access paths. Reduce 3 marks for each QEP.)



Late Penalty: Late assignments will be marked down **50% for each day it is late**. Any exceptions to this rule must be made prior to when the assignment is due and the excuse needs to be a good one - just too busy won't cut it. Individual exceptions are unfair to other students and hence they won't be made unless the circumstances are truly exceptional.

Plagiarism: All work submitted by you should be your own. **Copying or sharing of assignments would constitute plagiarism. Penalty will be given to those involved in plagiarism. This will be no exception once plagiarism is caught – being copied without notice won't excuse it; you should simply keep your assignment well.**

Submission: Please submit your assignment via BUMoodle