

# Relational Model

# Overview

## ❑ **ER Model** (last lecture)



- Conceptual design (High-level database design)
- *No direct relationship with database technologies*

## ❑ **Relational Data Model** (this chapter)

- Logical design
- Maps the *conceptual requirements* into the *data model* associated with a specific database management system.

# What is a Relation?

- ❑ A **relation** is a more concrete construction of the ER diagram
- ❑ A relation is (just!) a table!
- ❑ We will use **table** and **relation** interchangeably



sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

# What is a Relational Database?

❑ **Relational database:** a set of relations

❑ **Relation:** a named data table consisting of two parts:

➤ **Schema**


- A list of column/attribute names with their **data types**

**Students**(*sid*: char(20), *name*: char(20), *login*: char(10), *age*: integer, *gpa*: real)

➤ **Instance**

- Made up of zero or more tuples (or called *records*, *rows*)

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8



# What is a Relational Database?

❑ **Relation:** concept in relational data model:

➤ **Schema**

- specifies name of relation
- specifies name and type of each attribute

➤ **Instance**

- a table of tuples (or called *records, rows*)
- attributes (or called *fields, columns*).
- Number of tuples = cardinality
- Number of attributes = degree

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8


Cardinality = 3. and Degree = 5.

❑ **No two tuples are completely identical in a relation!**

# Relation vs. Relationship

## ❑ Relationship:

- concept in ER model
- Describes relationship between entities



## ❑ Relation:

- concept in relational data model
- Essentially a table (a set of tuples)

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

# The SQL Query Language

- ❑ A major strength of the relational model, which supports simple and powerful *querying* of data.
- ❑ Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
- ❑ SQL was developed by IBM in the 1970s.

# Create Relations in SQL

- ❑ Creates the **Students** relation.
  - Observe that the type of each attribute is specified.
  - When each tuple is added/modified, the DBMS must ensure that the tuple follows the type for each attribute.

```
CREATE TABLE Students
    (sid CHAR(20),
     name CHAR(20),
     login CHAR(10),
     age INTEGER,
     gpa REAL)
```



# Insert and Delete Tuples in SQL

- ❑ Can insert a single tuple using:

```
INSERT  
INTO Students (sid, name, login, age, gpa)  
VALUES ('53698', 'Bob', 'bob@comp', 18, 3.2)
```

- The INTO clause can be omitted if each value in VALUES corresponds to the correct attribute.
- Require single quotes for strings, e.g., 'Bob'.

- ❑ Can delete all tuples satisfying some condition(s) (e.g., name = 'Smith'):

```
DELETE  
FROM Students  
WHERE name = 'Smith'
```

# Data Types in SQL

❑ Different RDMS may support different data types

➤ **MySQL:** CREATE TABLE my\_table(date\_col **DATE**)

- E.g., ‘**2024-01-25**’ (YYYY-MM-DD)

➤ **PostgreSQL:** CREATE TABLE my\_table(date\_col **TIMESTAMP**)

- E.g., ‘**2024-01-25 10:23:54**’ (both date and time)

❑ MySQL as an example

➤ **String data types:** CHAR(size), fixed length string

➤ **Numeric data types:** INT(size), SMALLINT(size),  
FLOAT(size,d), FLOAT(p)

➤ **Date and time data types:** DATE, YEAR

# Creating Relations in SQL

- ❑ **char(n)**. Fixed length character string, with user-specified length n.
- ❑ **varchar(n)**. Variable length character strings, with user-specified maximum length n.
- ❑ **Integer(int)**. Integer (a finite subset of the integers that is machine-dependent).
- ❑ **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- ❑ **numeric(p,d)**. Fixed point number, with user-specified precision of p digits, with d digits to the right of decimal point. (ex., `numeric(3,1)`, allows 44.5 to be stored exactly, but not 444.5 or 0.32)
- ❑ **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- ❑ **float(n)**. Floating point number, with user-specified precision of at least n digits.

# Question 1

❑ Create the following relation using SQL.

- RK: integer
- GP: games played
- MPG: minutes per game

Points Per Game Leaders - Qualified

RK	PLAYER	TEAM	<u>GP</u>	<u>MPG</u>
1	Stephen Curry, PG	GS	79	34.2
2	James Harden, SG	HOU	82	38.1
3	Kevin Durant, SF	OKC	72	35.8
4	DeMarcus Cousins, C	SAC	65	34.6
5	LeBron James, SF	CLE	76	35.6
6	Damian Lillard, PG	POR	75	35.7
7	Anthony Davis, PF	NO	61	35.5
8	Russell Westbrook, PG	OKC	80	34.4

# Question 2

- ❑ Create the following Stocks relation using SQL, where Symbol, Name, Last, Change, and % Chg are represented by 20 characters, 100 characters, real number, real number, and real number, respectively.

Figure 3.31 Stock Quotations

Symbol	Name	Last	Change	% Chg
\$COMPX	Nasdaq Combined Composite Index	1,400.74 ▼	-4.87	-0.35%
\$INDU	Dow Jones Industrial Average Index	9,255.10 ▼	-19.80	-0.21%
\$INX	S&P 500 INDEX	971.14 ▼	-5.84	-0.60%
ALTR	Altera Corporation	13.45 ▼	-0.450	-3.24%
AMZN	Amazon.com, Inc.	15.62 ▲	+0.680	+4.55%
CSCO	Cisco Systems, Inc.	13.39 ▼	-0.280	-2.05%
DELL	Dell Computer Corporation	24.58 ▼	-0.170	-0.69%
ENGX	Enterprise Growth C	14.60 ▼	-0.210	-1.42%
INTC	Intel Corporation	18.12 ▼	-0.380	-2.05%
JNJ	Johnson & Johnson	53.29 ▼	-0.290	-0.54%
KO	Coca-Cola Company	56.70 ▼	-0.580	-1.01%
MSFT	Microsoft Corporation	53.96 ▲	+1.040	+1.97%
NKE	NIKE, Inc.	57.34 ▲	+0.580	+1.02%

- ❑ Perform the following actions in SQL:
  - a) Add a new record ('Google', 'Google Inc.', 450.00, 10%, 200.00).
  - b) Delete the records with *Last* < 15.

# Integrity Constraints

- ❑ **Integrity Constraints (ICs):** conditions that must be true for *any* instance of the database; e.g., **data type**.
  - ICs are specified when schema is defined.
  - ICs are checked when relations are modified.
- ❑ ICs are based upon the semantics of the application that is being described in the database relations.
- ❑ A **legal** instance of a relation is one that satisfies all specified ICs.
  - DBMS should not allow illegal instances.
  - Avoids data entry errors, too!

# Some Representative Examples of Integrity Constraints

- ❑ Primary key constraint
- ❑ Unique constraint
- ❑ Not NULL constraint
- ❑ Referential integrity constraint

# Primary Key Constraint



- ❑ A set of attribute(s) is a key for a relation if :
  1. No two tuples can have same value(s) in all the attribute(s) of the key.
  2. This is not true for any subset of the key (*minimal* requirement).
  
- ❑ Example:

Students(*sid*: char(20), *name*: char(20), *login*: char(10), *age*: integer, *gpa*: real).
  
- ❑ *sid* is a key. (What about *name*?)
  
- ❑ {*sid*, *gpa*} is not a key (Condition 2 is violated).



# Unique Constraint

❑ If there is more than one key for a relation (i.e., candidate keys), one is chosen as the primary key.

❑ Candidate key is unique.

❑ Example:

Students(*sid*: char(20), *name*: char(20), *login*: char(10), *age*: integer, *gpa*: real).

❑ *login* is a candidate key, which must be unique.

# Primary and Unique Constraints in SQL

- ❑ Primary key is specified using **PRIMARY KEY**.
- ❑ Candidate key(s) is (are) specified using **UNIQUE**.

```
CREATE TABLE Students  
    (sid CHAR(20),  
     name CHAR(20),  
     login CHAR(10),  
     age INTEGER,  
     gpa REAL,  
     PRIMARY KEY (sid),  
     UNIQUE (login) )
```

# Question 3

- ❑ Suppose that the table “Students” is created by the following SQL

```
CREATE TABLE Students
    (sid CHAR(20),
     name CHAR(20),
     login CHAR(10),
     age INTEGER,
     gpa REAL,
     PRIMARY KEY (sid),
     UNIQUE (login) )
```

- ❑ What’s the result of executing the following statements?
- INSERT INTO Students VALUES (‘00001’, ‘Bob’, ‘bob@comp’, 18, 3.2)
  - INSERT INTO Students VALUES (‘00001’, ‘Tom’, ‘tom@comp’, 18, 3.2)
  - INSERT INTO Students VALUES (‘00002’, ‘Bob’, ‘bob@comp’, 18, 3.2)

# Not NULL constraint

- ❑ If a column of a table is NOT NULL, you must give a value when you insert a tuple to the table.

- ❑ SQL Example:

```
CREATE TABLE Students
    (sid CHAR(20),
     name CHAR(20) NOT NULL,
     login CHAR(10),
     age INTEGER,
     gpa REAL)
```

- ❑ Remark: All attributes of the primary key MUST NOT be NULL.

# Referential Integrity Constraint

## ❑ Example: The Students and Enrolled Relations

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8


Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

- ❑ Assume that we want to insert a tuple ('50000', 'CS160', 'A') into Enrolled.
- ❑ Before we do so, we may want to make sure there exists a student in Students with *sid* = '50000'.
- ❑ This is called as **referential integrity**. (How to achieve this?)

# Foreign Key

- ❑ **Foreign key**: Set of attribute(s) in one relation that is used to 'refer' to a tuple in another relation (can be itself).
  - Must correspond to **primary key** of the referenced relation. Like a 'logical pointer'.
  - E.g., *sid* of **Enrolled** is a foreign key referring to **Students**.



```
Students(sid: char(20), name: char(20), login: char(10), age: integer, gpa: real)
Enrolled(sid: char(20), cid: char(20), grade: char(2))
```

- ❑ Can achieve referential integrity: make sure every sid in Enrolled exists in Students.

# Foreign Key in SQL

- ❑ Only students listed in the *Students* relation should be allowed to **enroll** for courses.
- ❑ But some tuples in *Students* may not be referenced.

```
CREATE TABLE Enrolled
(sid CHAR(20), cid CHAR(20), grade CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid) REFERENCES Students (sid) )
```



When a tuple is inserted in table *Enrolled*, what constraints are being checked?

**Enrolled**

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B


**Students**

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

# Foreign Key in SQL

❑ Attribute names can be different.

Students(sid: char(20), name: char(20), login: char(10), age: integer, gpa: real)  
Enrolled2(stuid: char(20), cid: char(20), grade: char(2))



```
CREATE TABLE Enrolled2
  (stuid CHAR(20), cid CHAR(20), grade CHAR(2),
   PRIMARY KEY (stuid, cid),
   FOREIGN KEY (stuid) REFERENCES Students (sid) )
```



# Foreign Key in SQL

❑ Foreign keys can refer to the same relation.

❑ Example:



**Students2(sid:char(20), name:char(20), login:char(10), age:integer, gpa:real, partner:char(20))**

```
CREATE TABLE Students2
  (sid CHAR(20), name CHAR(20), login CHAR(10),
   age INTEGER, gpa REAL, partner CHAR(20),
   PRIMARY KEY (sid),
   FOREIGN KEY (partner) REFERENCES Students2 (sid) )
```

- If a student has no partner, this attribute can be **NULL** (a special keyword in SQL denoting '*unknown*' or '*inapplicable*').
- NULL is allowed in non-primary keys, including foreign keys.

# Enforcing Referential Integrity via Foreign Key

## ❑ Example:

- Students and Enrolled Relations

- sid in Enrolled is a foreign key that references Students.

❑ What should be done if an Enrolled tuple with a non-existent student id is inserted? **Reject it!**

❑ What should be done if a Students tuple is deleted?

- Disallow deletion/update of a Students tuple that is referred to.  
(NO ACTION)

- Also delete/update all Enrolled tuples that refer to it.  
(CASCADE)

❑ Similar if primary key of Students tuple is updated.

# Enforcing Referential Integrity in SQL

- ❑ NO ACTION (*delete/update is rejected*).
- ❑ CASCADE (also delete/update all tuples that refer to deleted tuple).
- ❑ The DBMS adopts the default referential integrity if there is no specification.

```
CREATE TABLE Enrolled
(sid CHAR(20) DEFAULT '53688',
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students (sid)
ON DELETE CASCADE
ON UPDATE NO ACTION)
```

```
CREATE TABLE Enrolled
(sid CHAR(20) DEFAULT '53688',
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students (sid)
ON DELETE CASCADE
ON UPDATE CASCADE)
```

# Question 4

Consider the instances of Students and Enrolled relations on Slide 23. Suppose that the relation “Enrolled” is created based on the following SQL queries.

```
CREATE TABLE Enrolled
(sid CHAR(20) DEFAULT '53688',
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students (sid)
ON DELETE NO ACTION
ON UPDATE NO ACTION)
```

What happens if we perform the following operations for the instance “Students”:

- a) Delete the tuple with sid = 53666
- b) Insert a tuple with sid = 53600, name = ‘Edison’, login = ‘edison@cs’, age = 31, and gpa = 4.0.
- c) Update the tuple with sid = 53650 to 53700

# Question 5

- ❑ Repeat Question 3 if the relation “Enrolled” is created based on the following SQL queries.

```
CREATE TABLE Enrolled
(sid CHAR(20) DEFAULT '53688',
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students (sid)
ON DELETE CASCADE
ON UPDATE CASCADE)
```


# Summary

- ❑ Mostly commonly used table constraints
  - Domain (data type) constraint
  - Primary Key constraint
  - Unique constraint
  - Not NULL constraint
  - Referential Integrity constraint
  
- ❑ More general constraints can also be supported

# ER Model to Relation?

- ❑ How can we convert the high-level ER model to relational tables?
- ❑ Next we will explain the conversion for each component in an ER diagram.


# Entity Set to Relation



```
CREATE TABLE Employees
(ssn CHAR(11),
name CHAR(20),
lot INTEGER,
PRIMARY KEY (ssn))
```

### Employees

ssn	name	lot
R123	Alice	32
P625	Bob	87
A252	Candy	16




```
CREATE TABLE Departments
(did INTEGER,
dname CHAR(20),
budget REAL,
PRIMARY KEY (did))
```

### Departments


did	dname	budget
1	COMP	15000
2	MATH	15000
3	PHYS	12000




# Review: Cardinality Constraints




**1-to-1**



**1-to Many**




**Many-to-1**



**Many-to-Many**

*Translation to  
relational model?*

# m-m Relationship Set to Relation



**Works\_In**

<u>ssn</u>	<u>did</u>	since
R123	1	2011
R123	2	2016
P625	1	2018

❑ To convert a relationship set to a relation, it must include:


- keys for each participating entity set (as foreign keys)
- all descriptive attributes of the relationship

❑ What is the primary key of this relation?

❑ Which relations do we have for this ER diagram?

```
CREATE TABLE Works_In(  
    ssn CHAR(11),  
    did INTEGER,  
    since INTEGER,  
    PRIMARY KEY (ssn, did),  
    FOREIGN KEY (ssn)  
        REFERENCES Employees (ssn),  
    FOREIGN KEY (did)  
        REFERENCES Departments (did))
```

# 1-m Relationship Set to Relation



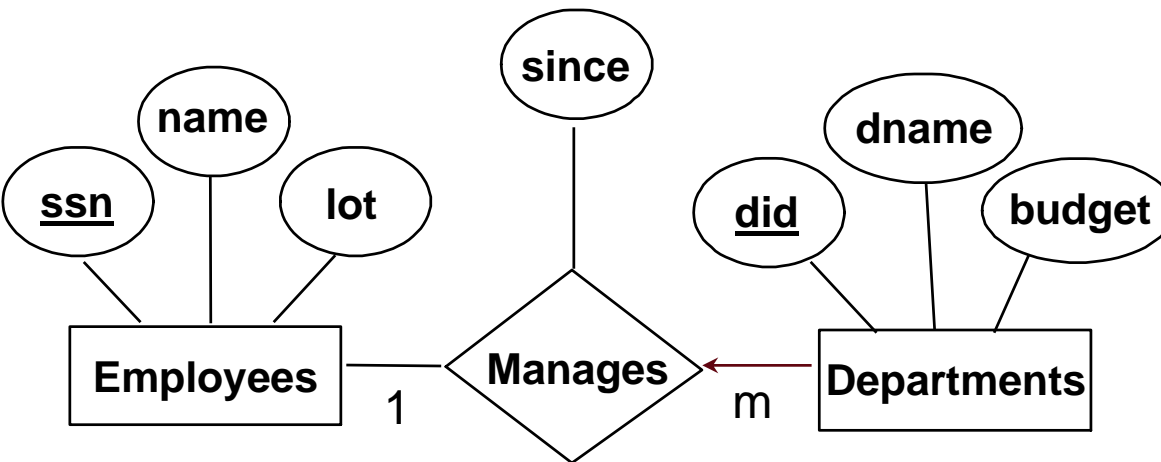
**Manages**

ssn	<u>did</u>	since
R123	1	2011
R123	2	2016
P625	1	2018

- ❑ Map the relationship “Manages” to a relation.
- ❑ What is the primary key of this relation?
- ❑ Which relations do we have for this ER diagram?

```
CREATE TABLE Manages(  
    ssn CHAR(11),  
    did INTEGER,  
    since INTEGER,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn)  
        REFERENCES Employees,  
    FOREIGN KEY (did)  
        REFERENCES Departments)
```

# 1-m Relationship Set to Relation



**Dept\_Mgr**


<u>did</u>	dname	budget	ssn	since
1	COMP	15000	R123	2011
2	MATH	15000	R123	2016

❑ Since each department has a unique manager, we can combine **Manages** and **Departments** together into one relation, called **Dept\_Mgr**.

❑ There are two relations, which are **Employees** and **Dept\_Mgr**.

```
CREATE TABLE Dept_Mgr(
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  ssn CHAR(11),
  since INTEGER,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn)
    REFERENCES Employees)
```

# 1-m Relationship Set (with Participation Constraint) to Relation




Dept_Mgr				
<u>did</u>	dname	budget	ssn	since
1	COMP	15000	R123	2011
2	MATH	15000	R123	2016

- Setting ssn to be NOT NULL makes sure every did (or department) in the Dept\_Mgr must have the corresponding ssn (or employee) to manage it. Therefore, every department must participate in this relationship.

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11) NOT NULL,  
  since INTEGER,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn)  
    REFERENCES Employees)
```

# 1-m Relationship Set (with Participation Constraint) to Relation



**Manages**

ssn	<u>did</u>	since
R123	1	2011
R123	2	2016


❑ Cannot capture participation constraints by adding the NOT NULL constraint to ssn.

❑ Because we can add many new tuples into the Departments relation, i.e., it cannot guarantee every did of Department appears in Manages.

```
CREATE TABLE Manages(  
  ssn CHAR(11) NOT NULL,  
  did INTEGER,  
  since INTEGER,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn)  
    REFERENCES Employees,  
  FOREIGN KEY (did)  
    REFERENCES Departments)  
(doesn't work)
```

# Hard to Capture All Participation Constraints!


□ Example:



- It is **insufficient** to use only PRIMARY KEY, UNIQUENESS, NOT NULL, and FOREIGN KEY to implement all participation constraints.
- More powerful features of SQL are needed to implement the relationship *Works\_In*.

## Question 6

Consider the following ER diagram.



Write the SQL query that can capture the cardinality constraint of the above ER diagram.



# Solution to Question 1

❑ Create the following relation using SQL.

➤ CREATE TABLE **Points Per Game  
Leader - Qualified**

(RK **INTEGER**,  
PLAYER **VARCHAR(20)**,  
TEAM **CHAR(10)**,  
GP **INTEGER**,  
MPG **REAL**)

Points Per Game Leaders - Qualified

RK	PLAYER	TEAM	GP	MPG
1	Stephen Curry, PG	GS	79	34.2
2	James Harden, SG	HOU	82	38.1
3	Kevin Durant, SF	OKC	72	35.8
4	DeMarcus Cousins, C	SAC	65	34.6
5	LeBron James, SF	CLE	76	35.6
6	Damian Lillard, PG	POR	75	35.7
7	Anthony Davis, PF	NO	61	35.5
8	Russell Westbrook, PG	OKC	80	34.4

# Solution to Question 2

❑ CREATE TABLE **Stocks**

(Symbol **CHAR(20)**,  
Name **CHAR(100)**,  
Last **REAL**,  
Change **REAL**,  
%Chg **REAL**)

❑ INSERT INTO Stocks(Symbol, Name, Last, Change, %Chg)  
VALUES('Google', 'Google Inc.', 450.00, 10%, 200.00)

❑ DELETE FROM Stocks WHERE Last<15

# Solution to Question 3

- ❑ The first record is successfully inserted into the table

sid	name	login	age	gpa
00001	Bob	bob@comp	18	3.2

- ❑ The DBMS rejects the 2<sup>nd</sup> statement since its sid=00001, which is the primary key and a record with sid=00001 already exists in the table.
- ❑ The DBMS rejects the 3<sup>rd</sup> statement since its login='bob@comp', which should be unique and a record with login='bob@comp' already exists in the table.

# Solution to Question 4

- ❑ The DBMS rejects this operation as there are three records with the foreign key sid=53666 in the Enrolled relation. The 'NO ACTION' referential integrity has been imposed to this foreign key.
- ❑ The DBMS executes this operation successfully.
- ❑ The DBMS rejects this operation as there is one record with foreign key sid=53650 in the Enrolled relation. The 'NO ACTION' referential integrity has been imposed to this foreign key.

# Solution to Question 5

❑ The DBMS executes this operation successfully.

Students

sid	name	login	age	gpa
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❑ The DBMS executes this operation successfully.

sid	name	login	age	gpa
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8
53600	Edison	edison@cs	31	4.0

❑ The DBMS executes this operation successfully.

sid	name	login	age	gpa
53688	Smith	smith@eecs	18	3.2
53700	Smith	smith@math	19	3.8
53600	Edison	edison@cs	31	4.0

# Solution to Question 6

```
CREATE TABLE Bosses
    (ssn    INTEGER,
     name   CHAR(20),
     PRIMARY KEY(ssn)
    );
```

```
CREATE TABLE Company
    (cid    INTEGER,
     cname  CHAR(20),
     property REAL,
     PRIMARY KEY(cid)
    );
```

```
CREATE TABLE Owns
    (ssn    INTEGER,
     cid    INTEGER,
     since   INTEGER,
     PRIMARY KEY (ssn, cid),
     FOREIGN KEY (ssn) REFERENCES
                                     Bosses(ssn),
     FOREIGN KEY (cid) REFERENCES
                                     Company(cid)
    );
```

# Conclusion

## ❑ Relational database

- Schema
- Instance

## ❑ Integrity constraints

- Primary key constraint
- Unique constraint
- Not NULL constraint
- Referential integrity constraint (foreign key)

## ❑ Mapping ER-diagram to Relational Models

- Handle the cardinality constraints in the ER-diagram.
- Handle the participation constraints in the ER-diagram.

## ❑ Basic SQL